

```
In [1]: ▶ import pandas as pd
import numpy as np
import sklearn as sk
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
```

```
In [2]: ▶ df=pd.read_csv("C:\\Users\\asus\\Desktop\\Boston.csv")
df
```

Out[2]:

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	blk
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	502	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391
502	503	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396
503	504	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396
504	505	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393
505	506	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396

506 rows × 15 columns



## Data Exploration

```
In [3]: ▶ df.shape
```

Out[3]: (506, 15)

In [4]: `df.describe()`

Out[4]:

	Unnamed: 0	crim	zn	indus	chas	nox	rm	
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	5
<b>mean</b>	253.500000	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	
<b>std</b>	146.213884	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	
<b>min</b>	1.000000	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
<b>25%</b>	127.250000	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
<b>50%</b>	253.500000	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
<b>75%</b>	379.750000	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	
<b>max</b>	506.000000	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	1

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   506 non-null    int64
1   crim        506 non-null    float64
2   zn          506 non-null    float64
3   indus       506 non-null    float64
4   chas        506 non-null    int64
5   nox         506 non-null    float64
6   rm          506 non-null    float64
7   age         506 non-null    float64
8   dis         506 non-null    float64
9   rad         506 non-null    int64
10  tax         506 non-null    int64
11  ptratio     506 non-null    float64
12  black       506 non-null    float64
13  lstat       506 non-null    float64
14  medv        506 non-null    float64
dtypes: float64(11), int64(4)
memory usage: 59.4 KB
```

In [6]: `df.isnull().sum().sum()`

Out[6]: 0

```
In [7]: df['chas'].value_counts()
```

```
Out[7]: 0    471  
        1     35  
        Name: chas, dtype: int64
```

```
In [8]: df['zn'].value_counts()
```

```
Out[8]: 0.0    372  
        20.0    21  
        80.0    15  
        22.0    10  
        12.5    10  
        25.0    10  
        40.0     7  
        45.0     6  
        30.0     6  
        90.0     5  
        95.0     4  
        60.0     4  
        21.0     4  
        33.0     4  
        55.0     3  
        70.0     3  
        34.0     3  
        52.5     3  
        35.0     3  
        28.0     3  
        75.0     3  
        82.5     2  
        85.0     2  
        17.5     1  
        100.0    1  
        18.0     1  
        Name: zn, dtype: int64
```

In [9]:

df.dtypes


```
Out[9]: Unnamed: 0      int64
      crim      float64
      zn      float64
      indus      float64
      chas      int64
      nox      float64
      rm      float64
      age      float64
      dis      float64
      rad      int64
      tax      int64
      ptratio      float64
      black      float64
      lstat      float64
      medv      float64
      dtype: object
```

In [10]:

```
df.corr()
plt.figure(figsize=(15,8))
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
In [11]: y3=df['medv']  
for i in df.columns:  
    x3=df[i]  
    plt.scatter(x3,y3)  
    plt.show()
```



## Data Preprocessing

```
In [12]: x_train=df.drop(['medv'],axis=1)
y_train=df['medv']
x_train
```

Out[12]:

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	blk
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	502	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391
502	503	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396
503	504	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396
504	505	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393
505	506	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396

506 rows × 14 columns



## Data Modeling

```
In [13]: from sklearn.linear_model import LinearRegression
```

```
In [14]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: df2=pd.read_csv("C:\\Users\\asus\\Desktop\\Boston_Test (1).csv")
df2.head()
```

Out[15]:

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black
0	351	0.07950	60.0	1.69	0	0.411	6.579	35.9	10.7103	4	411	18.3	370.7
1	352	0.07244	60.0	1.69	0	0.411	5.884	18.5	10.7103	4	411	18.3	392.5
2	353	0.01709	90.0	2.02	0	0.410	6.728	36.1	12.1265	5	187	17.0	384.7
3	354	0.04301	80.0	1.91	0	0.413	5.663	21.9	10.5857	4	334	22.0	382.8
4	355	0.10659	80.0	1.91	0	0.413	5.936	19.5	10.5857	4	334	22.0	376.0


```
In [16]: x_test=df2.drop(['medv'],axis=1)
y_test=df2['medv']
x_test
```

Out[16]:

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b
0	351	0.07950	60.0	1.69	0	0.411	6.579	35.9	10.7103	4	411	18.3	37
1	352	0.07244	60.0	1.69	0	0.411	5.884	18.5	10.7103	4	411	18.3	39
2	353	0.01709	90.0	2.02	0	0.410	6.728	36.1	12.1265	5	187	17.0	38
3	354	0.04301	80.0	1.91	0	0.413	5.663	21.9	10.5857	4	334	22.0	38
4	355	0.10659	80.0	1.91	0	0.413	5.936	19.5	10.5857	4	334	22.0	37
...	...	...	...	...	...	...	...	...	...	...	...	...	...
150	501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	39
151	502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	39
152	503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	39
153	504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	39
154	505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	39

155 rows × 14 columns

```
In [17]: ▶ y_pred=lr.predict(x_test)
plt.scatter(y_test,y_pred)
plt.show()
```



```
In [18]: ▶ from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score
mse=mean_squared_error(y_test,y_pred)
ae=mean_absolute_error(y_test,y_pred)
re=r2_score(y_test,y_pred)
print(mse)
print(ae)
print(re)
```

```
37.13832013971392
4.1529745209205755
0.44226910905239236
```

Score By Using Linear Regression

```
In [19]: ▶ lr.score(x_test,y_test)
```

```
Out[19]: 0.44226910905239236
```


## Random Forest



```
In [24]: from sklearn.ensemble import RandomForestRegressor  
regressor = RandomForestRegressor(n_estimators = 100, random_state = 0)  
regressor.fit(x_train, y_train)
```

Out[24]: RandomForestRegressor(random\_state=0)

```
In [21]: y_pred=lr.predict(x_test)  
plt.scatter(y_test,y_pred)  
plt.show()
```



```
In [22]: regressor.score(x_test,y_test)
```

Out[22]: 0.9574974486719926

```
In [ ]:
```